

Review of the `spike/config` Branch of `giv`

Overview & Architecture

The `giv` CLI is structured as a collection of POSIX shell scripts that break the functionality into modular components (e.g. commands, project metadata, history, LLM interaction, etc.). In the `spike/config` branch, we see a clear effort to simplify and refactor the codebase for maintainability and clarity. The core design goals are to keep everything POSIX-compliant (runnable in `/bin/sh` on Bash, Dash, Zsh, Ash, etc.), minimize external dependencies (only Git and standard tools like `grep`, `sed`, `jq`, etc.), and make the tool feel natural for Git users. Overall, the direction is good – splitting logic into dedicated scripts (e.g. `src/commands/...`, `project_metadata.sh`, `history.sh`, `llm.sh`, etc.) improves clarity, and the code largely adheres to shell best practices (quoting variables, using traps for cleanup, etc.). Below, I'll highlight specific issues and improvement opportunities in the code, focusing on redundancy, simplicity, and maintainability, while preserving functionality.

Removal of Redundant Code

Provider Scripts: A major improvement in this branch is the removal of specialized “provider” scripts in favor of a unified metadata handling function. Previously, separate scripts (e.g. `provider_python_toml.sh`, `provider_node_pkg.sh`, etc.) parsed project version info for different ecosystems (Python, Node, etc.). In `spike/config`, those have been deprecated and consolidated. Commit `f1bd745` explicitly notes “*Removed the `provider_python_toml.sh` script as it is no longer needed*” and “*Consolidated metadata extraction logic into `project_metadata.sh` for better maintainability*” ¹. This is a great change – it deletes a lot of duplicate logic (over **1200** lines removed ² ³) and replaces it with a single source of truth for version/title extraction. The new centralized function `get_metadata_value` in `project_metadata.sh` detects the project type and retrieves metadata like version or project name in one place ⁴. This eliminates redundant code paths and ensures all commands use the same logic for fetching project metadata, reducing bugs and maintenance effort.

Recommendation: Continue this direction by removing any remaining vestiges of the old provider system. For example, `provider_auto.sh` and `provider_custom.sh` are still present in the file list ³. If these are no longer used (likely superseded by `get_metadata_value` logic), they should be deleted to avoid confusion. Similarly, any global variables or config flags related to “providers” can be removed. For instance, if there was an environment variable to choose a provider, it can be retired in favor of auto-detection and the `--version-file/--version-pattern` options. This keeps configuration surface minimal and code paths straightforward.

Glow Integration: The branch introduces a new `dependencies/glow.sh` for installing the Glow CLI tool ⁵. Glow is a Markdown renderer for terminal. While it can enhance output display, it's an **optional** nice-to-have that adds complexity. The installation script and logic to invoke Glow (likely in `markdown.sh` or similar) should be isolated so that core functionality doesn't depend on it. Ideally, check for `glow` presence and use it only when available or explicitly requested. Automatic installation should be done cautiously –

e.g., ask for user confirmation, handle failures gracefully, and avoid doing anything on systems where it's not supported. Since one goal is "no dependencies other than Git or common tools," making Glow purely optional (or providing instructions for manual install) would align better with that goal. The code currently adds a script to manage Glow installation ¹; ensure this doesn't run without user intent. Simplicity is key – we don't want to complicate the user's environment unexpectedly.

Config File Handling (`giv config`)

The branch adds a **config subcommand** and greatly improves how configuration is managed. Notably, `giv config` now behaves much like `git config`, which will feel familiar to users. You support listing all settings, getting, setting, and unsetting keys, and even a shorthand for get/set without explicit flags (e.g. `giv config user.name "Alice"` to set, or `giv config user.name` to get) ⁶ ⁷. This is excellent for usability. A few points to consider:

- **File Location:** By default, config is stored in `~/.giv/config` (with `$GIV_HOME` defaulting to `~/.giv`). The tests use a fake home to verify this. This is fine, but consider also supporting the XDG Base Directory standard (e.g. `$XDG_CONFIG_HOME/giv/config`) if present. Many modern CLI tools prefer `~/.config/toolname/...`. You could check `XDG_CONFIG_HOME` and fall back to `HOME/.giv`. This makes it more flexible for users who organize dotfiles. It's minor, but aligns with best practices.
- **Project-Specific Config:** It's not clear if you plan to allow project-local config (e.g. a `.giv/config` in the repo). Currently, `--config-file` lets the user specify any config path manually, and the branch introduced a `--config-file PATH` global option. If it's not already implemented, consider automatically looking for a config in the Git repository root (e.g. `.giv/config` in the repo) and loading that **in addition to** the global config. This way, a project can carry specific settings (like its version file pattern, etc.) under version control. The commit message doesn't explicitly mention project config, but adding that would enhance flexibility.
- **No-Argument Behavior:** The usage string suggests `giv config` with no further args should perhaps list all config (`--list`) or show help ⁶. In Git, `git config -l` lists, and `git config [name]` gets a value. In your implementation, I see that you handle `list`, `get`, `set`, etc., and also a `*` catch-all for the shorthand. Ensure that `giv config` **with no arguments** produces a useful result – either treat it as `giv config list` (preferred) or print a help/usage message. As of the latest diff, an empty `$1` would fall through the `case` without matching any pattern (since you match `*-|help` for unknown and have a `*`) for shorthand). You did add a usage output in the `help` pattern, but consider explicitly handling no arguments. For example, you could default `$cmd="--list"` if `$# -eq 0` before entering the case. This would align with user expectation that running the config command alone lists all settings.
- **Robustness & Error Handling:** The config script improvements are well done. You now guard every operation with proper checks:
 - If the config file doesn't exist, `--get`, `--unset`, etc. will print "config file not found" and return error ⁸ ⁹.

- You *no longer always create an empty config file on every run*, which is good – originally the code did `mkdir -p && touch` on each `giv config` invocation, which was unnecessary. Now you only create the file when setting a value (if it doesn't exist) ¹⁰ ¹¹. This avoids clutter (no empty file unless needed).
- The **malformed line check** was updated to ignore blank lines and comments ¹². Great move – now users can add comments (`# ...`) or empty spacing in their config for readability without causing “Malformed config” errors. This check now only flags lines that are not blank/comment and not of the form `key=value` ¹².
- **Environment Overrides:** A valuable addition is that environment variables override config values. For example, `giv config --get foo.bar` will first check if an env var `GIV_FOO_BAR` is set and return that ¹³ ¹⁴. If not, it then falls back to the config file ¹⁵ ¹⁶. This means users can temporarily override config by exporting an env var, which is intuitive and useful. Your implementation uses `eval` to fetch the dynamic env var, which is fine (since the var name is constructed safely). Just be sure to quote properly (which you did with `eval "env_val=\"\${$env_key-}\""` ¹⁵). Also, when listing config (`--list`), you convert any `GIV_FOO_BAR=value` lines in the file to a user-friendly `foo.bar=value` format ¹⁷ ¹⁸. This is clever – it hides the implementation detail of using `GIV_` prefixes internally, showing config in a clean, dotted-key format. It might be worth also indicating (maybe via a comment in `giv config --list` output or docs) that certain values might be coming from environment. For instance, if an env override is in effect, `--get` returns the env value even if the file has something different. This could confuse users viewing the file. Perhaps a note in documentation or an indicator in `--list` output for env-sourced values would help. However, that might complicate output, so it's optional.
- **Key/Value Handling:** A few small suggestions to simplify and harden:
 - **Escaping regex:** When using `grep -E "^$key="`, if `key` contains regex special characters, this could misbehave. In practice your config keys likely only contain alphanumeric, dots, and underscores (by convention), so it's mostly fine. Dots *are* special in regex (matching any char). For example, a key `foo.bar` will match `fooZbar` in the grep because the `.` is not escaped. To be safe, you could escape `.` in the pattern. A quick way is using grep's fixed string option: `grep -F -m1 "$key=" "$GIV_CONFIG_FILE" | cut -d'=' -f2-` (with `-m1` to stop at first match). Since the file format is simple and only one occurrence of a key is expected, using fixed-string matching avoids regex pitfalls entirely. Alternatively, use a small function to escape regex metacharacters in `$key`. This will prevent surprising behavior if someone uses, say, a key with a dot or `+` in it.
 - **Unique keys:** The config file shouldn't contain duplicate keys, and your `--set` command enforces that by removing any existing lines for the key (and its `GIV_` form) before writing the new value ¹⁹ ²⁰. Good. If multiple duplicates somehow existed, your `grep | cut` would return multiple lines. It might be prudent to handle that scenario with an additional check (e.g. if `val` contains a newline, perhaps warn that the config had duplicate entries). But since your `set` and `unset` commands clean duplicates, this should not occur unless the user manually edited the file incorrectly. Not a big issue.
 - **Use of `echo` vs `printf`:** In a few places, you use `echo` to output potentially arbitrary strings. E.g., in `giv config --list`, you do `echo "$key=$value"` for normal lines ²¹. If a value begins with `-` or contains `\c` or other strange escape, `echo` can behave unpredictably (since its behavior with options is not consistent across shells). It's safer to use `printf '%s\n' "$line"` or

similar to ensure the output is exactly as intended. This is a general shell best practice: use `printf` for variable-driven output. The same goes for error messages: `echo "Unknown option: $cmd" >&2` is fine because `$cmd` in this context will start with `-` or be a word like "help", which won't confuse `echo` (and you redirect to `stderr`), but using `printf "%s\n" "$msg" >&2` is more robust. Consider auditing all `echo` calls and switching to `printf` where appropriate – especially if the content might be user-supplied. This will make the scripts more portable (some shells have `echo` that interprets `-e` sequences, etc.) and avoid odd edge cases.

- **Internal vs User Representation:** As noted, the config uses `GIV_FOO_BAR` variables under the hood and presents them as `foo.bar`. This is a smart approach to map to environment variables. The code converting between them is a bit repetitive (there are a few places where you do the `tr '[:lower:]' '[:upper:]'` and replace `.` \rightarrow `_`). This could be refactored into a tiny helper function, e.g. `normalize_key()` and `denormalize_key()`, to convert from user key to env var and vice versa. That would avoid repeating the `tr ... | tr ...` pipelines everywhere (currently in at least `get`, `set`, `unset` cases) and make it easier to adjust if needed. Even without using a function (since defining a function for a one-liner might be overkill in shell), you could do the transformation once and reuse it. For instance, in `--set|set`, you compute `givkey` and then again compute `writekey` in the same way ²² ²³. `writekey` is actually identical to `givkey` in content. You could just reuse `givkey` for writing, since you always store in `GIV_` form. Not a big issue, but avoids an extra process call and potential typo divergence. Similarly, in `--get`, you build `env_key` twice (once for checking env, once in the `grep` fallback) ¹⁵ ²⁴. One calculation of the uppercase/underscore key would suffice. These micro-optimizations make the code slightly shorter and clearer (one source of truth for that conversion logic).
- **Exit Codes:** The config functions appropriately return non-zero on failures (e.g. returning `1` when file not found or malformed config) to signal error. Ensure that these propagate out of the `giv` main script. If you're calling `giv_config` from within a case in the main command dispatcher, you may need to capture its return and `exit $?` or so. In a unified script, if the function returns non-zero, and you have `set -e`, that might terminate the script immediately (depending on context). Be mindful of how the `set -e` setting interacts with these returns. Usually, returning an error from a function in a case statement won't break the script if not explicitly checked. You might actually want to use `exit 1` in the function for fatal errors to ensure the whole CLI exits with error. Just double-check that the exit status of `giv config` is correctly reflecting success/failure. This is important for scripting use of `giv`.

Project Metadata & Version Extraction

As mentioned, the shift to a centralized metadata module is a significant improvement. The new `project_metadata.sh` (and possibly a `project/metadata.sh` it replaced) now handles retrieving things like project version, name/title, etc., by detecting the project type and looking in the appropriate files

⁴ .

Project Type Detection: It appears the script determines the project type during initialization (likely by seeing which files are present – e.g. if `package.json` exists, it's a Node project; if `pyproject.toml` exists, Python; etc.) and then stores that info (perhaps in a variable or config) for use by `get_metadata_value` ²⁵. This is a good approach: detect once, use many times. Ensure this detection is

robust and unambiguous. If a repo has multiple indicators (e.g. both a `package.json` and a `pyproject.toml` – not common, but possible in a multi-language monorepo), decide how to prioritize or allow the user to override. The old “provider_auto” logic may have handled such precedence; make sure the new function does too. For maintainability, consider using a simple case list: e.g.

```
if [ -f pyproject.toml ]; then project_type=python;
elif [ -f package.json ]; then project_type=node;
elif [ -f setup.py ]; then project_type=python;
# etc...
else project_type=unknown; fi
```

Storing `project_type` (maybe in a global like `GIV_PROJECT_TYPE`) is helpful, and I suspect you do that via the config or environment.

Version Extraction: The new approach presumably uses lightweight parsing (grep/sed) instead of heavy external tools. This aligns with the “no dependencies” goal. For example, reading `package.json` for version could be done either with `jq` (if allowed) or with grep. If you want zero dependencies, you might have implemented a grep-based JSON parse. That can be error-prone (JSON is a bit complex), but if you assume well-formatted package.json with `"version": "x.y.z"`, on one line, a regex can fetch it. Using `jq` would be more robust for JSON, but that introduces a dependency that not all systems have by default (though `jq` is commonly available, it's not as ubiquitous as `grep`). Since the policy is to allow “common things like jq if needed,” it might be acceptable to use `jq` here. I would lean towards using `jq` for reading JSON and maybe even TOML (if a `tomlq` or similar existed) for reliability. But if you choose not to depend on that, ensure the grep patterns are anchored and specific (to avoid matching other JSON keys that contain “version”). Perhaps something like:

```
grep -E '"version"[[:space:]]*:' package.json | head -1 | sed -E
's/.*"version"[[:space:]]*:[[:space:]]*"([^"]+)"*/\1/'
```

This gets the value of the first “version” key. It’s a bit hacky but works in simple cases. Just document that assumption.

For Python, removing `provider_python_toml.sh` suggests you no longer use Python’s `tomllib` or `pip` to get the version. Likely you grep `pyproject.toml` for a version field (PEP621 defines `version = "x.y.z"` in `[project]`). Again, make the grep strict: e.g. `grep -m1 -E '^version\s*=' pyproject.toml` and extract the quoted value. TOML allows single quotes or double quotes; ensure your regex covers both or use a minimal toml parser approach (since dependency-free, regex is fine, just be aware of edge cases like comments or multi-line values, which *shouldn't* happen for version).

Custom Version Sources: The presence of `provider_custom.sh` in the old system implies you allowed a custom command or file for version. If that’s still desired, you could incorporate it by letting users configure a shell command to get the version (perhaps via an environment variable like `GIV_VERSION_CMD`). In the new design, it might be even simpler: if the user provides `--version-file` or `--version-pattern`,

that essentially is a custom specification. So a separate custom provider script may be unnecessary. If any code related to `provider_custom` remains, consider replacing it with a direct check: e.g. *if* user provided a custom version file/pattern via CLI or config, use that to grep the version. Otherwise auto-detect project type and use defaults. That appears to be the approach now, which is good.

Redundancy: With the unified `get_metadata_value`, any functions like `get_project_version`, `get_project_title` etc., might be obsolete. Indeed, I see in the diff that calls to `get_project_version` were replaced by `get_metadata_value "version"` ²⁶. If `get_project_version` and similar older functions still exist in the code, you can remove them to avoid confusion. Make sure all references are updated to use the new unified function. The commit diff shows that was done in `commands.sh` for the `document` flow ²⁶ and likely elsewhere.

Error Handling: When metadata is missing or fails to parse, ensure the user gets a clear error. For example, if it's a Node project but `package.json` has no `"version"` field (some packages might omit it if not a published package), what happens? The script might return empty and continue. It might be better to detect an empty result and treat it as an error or at least warn. The commit shows instances of `get_metadata_value "version" "$commit" || { print_error "..."; return 1; }` in some contexts ²⁷, indicating you do handle failures gracefully in some places. Propagate such errors up so that, for example, `giv changelog` doesn't silently proceed with a blank version if none found – it could note “Version not found” or similar. Since one goal is to generate accurate release notes, knowing the version is important; failing clearly is better than guessing wrong.

Flexibility: The unified approach makes it easier to extend to new project types if needed (just add a case for that type in one place). This meets the goal of flexibility. Just be sure to keep it **simple** – avoid complex parsing. If a project doesn't fit a known pattern, instruct the user to use `--version-file` / `--version-pattern` (which you already support) rather than trying to handle every edge case automatically. This balances simplicity with flexibility.

Git History & Diff Collection

Since **giv** revolves around turning Git history into prompts, the parts of code that collect Git data are critical. It looks like you have a `history.sh` that provides utilities for summarizing commit histories, diffs, and extracting TODO comments ²⁸. A few things to consider in this area:

- **Unified History Logic:** Ensure that all subcommands that need to gather Git info use the same underlying functions. For example, generating a commit **message** for the working tree, a **summary** for a range, or **release notes** for two tags all require obtaining diffs, commit messages, or file lists. It's best if one module (`history.sh`) handles “given a revision range (or special flags like `--cached`), produce the set of changes and related metadata.” I suspect you have functions like `get_diff` or `collect_changes` in `history.sh`, and indeed commit notes mention “*summarizing Git history, extracting TODO changes, and caching summaries*” ²⁸. This is good for avoiding duplication. For example, if both `summary` and `changelog` need to list changed files, they should call the same function rather than each running separate `git diff` commands with slightly different flags.

- **Revision/Range Parsing:** The tool accepts any Git revision or range. That means the user might pass `HEAD~3..HEAD`, `v1.2.0..v1.3.0`, or even the special `--cached` or `--current` flags as pseudo-revisions. In a pure POSIX shell context, handling an argument like `--cached` is tricky because it looks like an option. It appears you treat `--cached` and `--current` as valid “revision” inputs rather than options. The README’s usage confirms this (they are listed as acceptable “revision” specifiers). In implementation, you likely have to differentiate between global options vs revision arguments. If you haven’t, I recommend doing argument parsing in two stages:

- Parse **known options** (like `-h`, `--help`, `--verbose`, `--dry-run`, `--config-file`, etc.) first. Remove them from `$@`.
- What remains can be interpreted as `[revision] [pathspec]` in that order (with 0, 1, or 2 positional arguments possibly left).

During this, treat `--cached` and `--current` specially: if `$1` is `--cached` or `--current` and wasn’t caught as a global option, then set a flag like `use_index=true` or `use_worktree=true` and consume that argument as the “revision”. For example:

```
case "$rev_arg" in
  --cached) rev_spec="--cached";;
  --current) rev_spec="--current";;
  *) rev_spec="$rev_arg";;
esac
```

Then inside your history functions, use those flags to decide the git commands. It sounds like you did implement this, but double-check edge cases. For instance, if the user passes *both* `--cached` and a range (not that they should, but if someone did `git summary --cached HEAD~5..HEAD`), your parser should probably error out or give precedence to one.

- **Working Tree vs Index vs Commits:** Clarify what `--current` means. From context, I suspect:
 - `--cached` = staged changes (index vs HEAD)
 - `--current` = all local changes (staged + unstaged vs HEAD)

Ensure the diffs you produce reflect that: - For `--cached`, you’d run `git diff --cached HEAD` (or just `git diff --cached` which implies HEAD). - For `--current`, you’d run `git diff HEAD` (which includes unstaged). If untracked files are present, `git diff` won’t show them. You might use `git diff --stat` or `git ls-files` for untracked if needed, but that might be beyond scope. Perhaps “current” just means “including unstaged changes that `git diff` can show”.

Also, when generating a commit **message** (`git message`), I suspect: - No revision arg => use `--current` by default (i.e. propose a message for all changes in the working tree that will be committed). - If `--cached` is given => only consider staged changes (so you can commit what’s staged, leaving unstaged for later). - If a specific revision (like a commit SHA) is given to `message`, perhaps you allow that to generate an AI message for an existing commit (though that’s an unusual case). Possibly `message` subcommand with a revision acts like summarizing that commit’s diff.

In any case, the code should handle these consistently across subcommands. It might help to have a single function like `prepare_diff_data(rev_spec, pathspec)` that sets up temp files or variables with the diff, list of commits, etc., according to the spec. This can be reused by `cmd_message`, `cmd_summary`, etc.

- **Performance and Large Histories:** Be mindful that a range like `v1.0.0..HEAD` could include hundreds of commits or a huge diff. The tool should handle it gracefully without excessive memory use. You generally avoid storing large outputs in variables (good), using files instead. For example, you likely write the `git diff` output to a temp file in `$GIV_TMPDIR` and maybe commit messages to another. This is the right approach. I see usage of `portable_mkdir` and `mkdir` in the code for creating temp files ²⁹ ³⁰. Make sure these files are cleaned up. You have a `GIV_TMPDIR_SAVE` debug option to preserve them (nice for debugging), otherwise you should remove them on exit. Set `trap 'rm -rf "$tmpdir"' EXIT` after creating your tempdir, for example. In the config example, `GIV_TMPDIR_SAVE` is documented ³¹; ensure the code actually implements it (likely in `system.sh` or `init.sh`).

Also, consider using `git --no-pager` and specific flags to limit output where appropriate. For instance, if generating a *summary*, you might not need the full diff of every commit, maybe just commit messages and short stats. But since you feed the LLM with diffs for context in some cases (especially for commit message generation), you probably do need unified diff content. Just be aware of size: maybe truncate extremely large diffs or warn if diff is huge (LLMs have context limits). This might be beyond the core code, but it's something to keep in mind for future improvements (like, maybe skip vendor/ or large binary diffs by default, unless included via `pathspec`).

- **TODO Extraction:** The tool scans for TODOs in code to provide context to the LLM about unfinished work. The `history.sh` mention of extracting TODO changes suggests you detect lines with “TODO” in the diff or in files. The implementation likely uses `grep -R` on the `pathspec` for the given pattern. Since the user can specify a regex via `--todo-pattern`, you pass that to `grep`. Make sure to quote it properly in the shell (likely need to use `grep -E` and `printf` or separate variable to avoid shell interpreting backslashes). Also, honor `--todo-files` (`pathspec` for which files to scan). If `--todo-files` is given, restrict `grep` to those files (maybe by using `git ls-files -z -- '$pathspec'` piped to `xargs grep`, or simpler, use `grep -R --include patterns`). Because this is user-configurable, just ensure if no TODO pattern is given, you either skip this step or use a sensible default pattern (the README example shows `TODO\w+`

`:` as an example pattern).

A potential improvement: If no `--todo-pattern` is provided but `--todo-files` is, you might still want to search for a default “TODO” string. Conversely, if a pattern is given but no `--todo-files`, you might search in all files by default (or all tracked files). Clarify these defaults in docs and code. Also, scanning the entire repo for “TODO” could be slow in big projects; if the user knows which files matter (e.g. only code files), they should specify it. Perhaps default to `--todo-files` = tracked code files only (exclude binaries, docs) if not provided, to avoid noise. In any case, your modular approach in `history.sh` likely encapsulates this well.

- **Git Command Errors:** Always check the return status of `git` commands. For example, if a user specifies an invalid revision range, `git diff` / `git log` will exit non-zero. With `set -e` that might terminate the script unexpectedly if not caught. It would be better to catch these and print a

friendly error like “Invalid revision range specified” and exit gracefully. You can do this by capturing output and status explicitly, or using `if ! git diff ... >file; then print_error "Git diff failed"; exit 1; fi`. The tests might cover some error scenarios (e.g. the `test_commands.bats` probably checks unknown subcommands, etc., but maybe not invalid revisions). It’s worth adding tests for a bad rev (like `giv summary badref`) to ensure you handle it. This will make the tool more robust.

LLM Integration (Local & Remote Models)

The app can call either a local model (via Ollama CLI) or a remote OpenAI-style API. In the code, this is likely handled by `llm.sh`. Key points:

- **Local vs Remote Selection:** You provide `--model-mode [auto|local|remote|none]` and related options like `--model` (for local model name), `--api-url`, `--api-model`, and environment variables `GIV_API_KEY`, etc. The code probably decides at runtime which mode to use. Ensure the logic for “auto” mode (probably: use local if Ollama is installed and `--model` given, otherwise fallback to remote) is clear and documented. If `--model-mode auto` and both local and remote are configured, which takes precedence? The expected behavior might be: if an `ollama` CLI is present and a `GIV_MODEL` is set (or default), use that; otherwise, if an API key is present use remote; if neither, error out or `--model-mode none` which presumably means skip the AI call (maybe for debug).
- **API Calls (Remote):** If using `curl` to call the OpenAI API (or compatible endpoint), make sure to handle errors and parse the JSON reliably. Using `jq` is highly recommended here to extract the `choices[0].message.content`. If `jq` is not available, you could attempt `grep`/`sed` on the JSON, but that’s very brittle. Since you explicitly allow `jq` as a common dependency, I would use it. It simplifies things:

```
response=$(curl -sS -H "Authorization: Bearer $GIV_API_KEY" -H "Content-Type: application/json" -d "$payload" "$GIV_API_URL")
content=$(printf "%s" "$response" | jq -r '.choices[0].message.content' 2>/dev/null)
```

Then check `jq`’s exit status and whether `content` is null or empty. If the API returned an error (HTTP error or JSON with error), handle that gracefully (maybe print the error message from JSON or at least notify the user something went wrong). Right now, it’s not clear if error handling for network/HTTP issues is implemented – it should be. For instance, if `curl` fails (no internet, or non-200 status), you might end up feeding an empty or error HTML into `jq`. You can detect `curl` failure by adding `--fail` to treat `HTTP ≥ 400` as error (then check `$?`). Or parse the HTTP status from the output if needed. Since this is a CLI likely used interactively, it may be acceptable to just print the error and not do elaborate retries.

- **Ollama (Local):** For local LLM via Ollama, presumably you run `ollama generate $modelName < prompt.txt`. The output is likely plain text. Confirm how you capture it. If the Ollama CLI streams tokens (like its normal interactive mode), you might want to pass `--json` or something to get a full

completion. (Ollama does have a JSON output mode as well). If you rely on plain output, just ensure to capture it fully. Also, check Ollama's exit codes – if it fails to find the model or errors, catch that and advise the user (e.g. “Local model not found. Install it or use --model-mode remote.”).

- **None mode:** If `--model-mode none` is used, I assume you skip calling any model and perhaps just print the prepared prompt to stdout or to the output file. This is a neat feature for debugging (seeing the prompt content). Ensure that when mode is “none”, the code path indeed bypasses the curl/ollama call, and that the output handling still writes something sensible (maybe the raw prompt into the output file or just does nothing). The documentation says “none” for model mode – perhaps that means “don’t do AI at all, just prepare outputs” which could equate to an identity transformation. It might be used to just assemble a changelog using only existing info (though the primary value of giv is the AI generation, “none” might just be for testing).
- **Streaming vs Wait:** Currently, it looks like the design is simple: call the model (which might take several seconds) and then write the result. There’s no streaming of partial output (which is fine for a first iteration). It keeps the code simpler. Just make sure to use `printf` or `tee` carefully if writing to both console and file. Likely you store the AI output in a temp file first, then decide to insert it into a changelog or print it. That approach is fine.
- **Security:** Since this is about using AI, be mindful of not exposing secrets inadvertently. For example, if the git diff contains secrets or if the user’s prompt template inadvertently includes something sensitive, it will be sent to the AI (remote). That’s more of a documentation concern. But one thing in code: you do source a config file (shell script) and `.env`. If those contain malicious code, sourcing them is a potential vector. However, given these are under the user’s control (their own config), it’s an acceptable risk; just document that “the config file is a shell script that will be executed”. Alternatively, if you wanted to be safer, you could parse the config file instead of sourcing, but that’s not trivial in POSIX sh (since you allow environment overrides, sourcing is the straightforward way). I think it’s fine – similar to how many CLIs source a profile script.

Prompt Templates & Markdown Processing

The tool uses prompt templates (in `templates/`) to structure what it sends to the LLM. It likely also does some token replacement (like inserting the collected diff, commit summaries, etc. into the template). The tests mention `replace_tokens.bats`, so you have a function to do this substitution.

Template Simplicity: Using plain Markdown files with placeholders is a good, simple approach. Common placeholders might be things like `{{CHANGES}}`, `{{VERSION}}`, `{{TODOS}}`, etc. The `replace_tokens.sh` (or perhaps within `markdown.sh`) likely implements something akin to:

```
sed -e "s/{{TOKEN}}/$replacement/g" ...
```

Be very careful with this if the replacement text can contain characters that are meaningful in sed replacement (especially `&` or backslashes). Ideally, escape the replacement text. If using GNU sed, you can use the `s|||` with `|` as delimiter and do something like:

```
escaped=$(printf '%s' "$replacement" | sed -e 's/[\\&]/\\&/g')
sed -e "s|{{TOKEN}}|$escaped|g" template.md
```

This replaces `/` and `&` with escapes to not break the sed. Also, ensure the placeholder itself does not contain regex special chars or at least escape them (the double braces might be interpreted as literal braces in BRE, which are not special unless forming a quantifier like `{2}`, but to be safe you can do `{{TOKEN}}` as a fixed string by turning off regex for those or using `grep -F` equivalent in sed by prefixing with `\Q` in GNU sed or simply constructing a sed script that treats braces literally).

An alternative approach that avoids these issues is to use shell parameter expansion in a here-doc. For example:

```
prompt=$(< "$template")    # read whole template
prompt=${prompt//\{\{CHANGES\}\}/"$changes"}
```

This will replace all occurrences of `{{CHANGES}}` with the content of `$changes`. In bash you could do that easily; in POSIX sh, parameter substitution with patterns is somewhat limited (it doesn't support multi-character substring replacement easily – though some shells support `${var//find/replace}` as an extension). So you likely opted for `sed`. Just test it with edge cases (like if the diff contains `&` or `\`, does it appear correctly in the final prompt?). Your `markdown.bats` test might cover some of this.

Glow vs Markdown: If Glow is installed, perhaps you plan to use it to display the final AI output (which is Markdown) in a nicely formatted way in the terminal. That's cool, but should be optional. Perhaps the `announcement` subcommand or others that produce Markdown could pipe through `glow` if user is in an interactive terminal. Implementation-wise, you might have a `view_markdown()` function in `markdown.sh` that does:

```
if command -v glow >/dev/null; then
    glow -p "$file" # preview in pager
else
    cat "$file"
fi
```

Something like that. That's a reasonable approach. Just ensure it doesn't interfere with non-interactive use (if output is being redirected, you might not want to use glow's pager). One approach is to detect if `stdout` is a TTY. For example, `if [-t 1] && [-x "$(command -v glow)"]; then ...`. If output is being piped or redirected, just output raw markdown. This respects common CLI behavior (never use pager/interactive output when not attached to TTY).

Replacing Tokens Not Needed: The new design might remove the need for a `build_prompt` script that was in earlier versions. If you simplified how prompts are constructed (maybe doing it inline in the subcommand logic), be sure you're not leaving behind any unused code. The commit list shows `build_prompts.bats` and `markdown.bats` tests ³², implying those functions still exist. If

`build_prompt()` and `generate_from_prompt()` are part of `commands.sh`, ensure they are minimal and not duplicating what could be a generic mechanism. Actually, given the branch includes a new `document.sh` command ³³ and mentions diagrams of how data flows ³⁴, it looks like you've thought through the architecture of combining various pieces. The `document` subcommand seems to orchestrate multiple modules to produce a full document (perhaps a combination of summary, changelog, etc.). This is beyond the scope of a simple review, but it indicates you have a high-level workflow in mind.

Maintaining Simplicity: While the architecture is now nicely modular, guard against over-engineering. Each script should have a clear, single responsibility and minimal inter-dependencies: - `system.sh` - low-level helpers (e.g. `load_env_file`, portable `mktemp`, debug printing, trap setup). - `project_metadata.sh` - determine and retrieve project meta (version, title). - `history.sh` - gather git diffs, commit logs, TODOs, possibly caching interim results. - `llm.sh` - handle calling the AI (local or remote). - `commands/*.sh` - implement each user-facing subcommand (using the above utilities). - `commands.sh` (if still present) - might coordinate complex flows or hold shared command logic. The diff snippet shows it's used for the `document` command logic and to define some global variables for linting ³⁵ ³⁶. Perhaps it used to contain all commands and now has been slimmed down. Ideally, all user subcommands could be in `commands/` directory as separate files, and `commands.sh` could just iterate through them or serve as a dispatcher. Actually, I see `giv.sh` in the file tree ³⁷ - likely **`giv.sh`** is the main entrypoint script that parses args and dispatches to the appropriate `cmd_*` function or script. That would be a central piece to keep simple and robust.

Familiarity to Git Users: The CLI is clearly modeled after git in many ways (subcommands like `summary`, `changelog` akin to `git log` outputs, the `config` command similar to `git config`, etc.). This is great. One more idea: consider supporting the `GIT_DIR` and `GIT_WORK_TREE` environment variables if not already. Git does, to operate in a non-standard repository location. If someone runs `giv` in a subdirectory of a repository or from outside the repo but with those env vars set, do your git commands respect that? Usually, if you just call `git` normally, it will pick up those env vars. So probably nothing needed except not interfering. Just avoid doing `cd` to the repo root unless necessary. If you *do* `cd $(git rev-parse --show-toplevel)` at any point (to make relative paths consistent, etc.), be careful to preserve the original cwd if needed. It might be fine to always operate at repo root; just note it in code comments.

ShellCheck and POSIX Compliance:

The branch includes a `.shellcheckrc` so you are linting the scripts, which is excellent. A few specific ShellCheck/POSIX points:

- **`local` in POSIX:** If you are targeting pure POSIX `sh`, technically `local` is not specified by POSIX. However, in practice, `local` is supported by **`bash`**, **`ksh`**, **`zsh`**, and even **`dash`** (Dash does accept `local` inside functions). It's not supported by *pure* `/bin/sh` on some older systems (like Solaris's `/bin/sh`). Since you list Dash and Ash (BusyBox) as supported, note: BusyBox **`ash`** **does** support `local`, and Dash does as well (they treat it as a no-op outside functions). So using `local` is generally safe on modern systems and ShellCheck (when set to `shell=dash`) will only warn (SC3045) if it's concerned about portability to a theoretical strictly POSIX shell. If you decided not to use `local` to avoid that warning, you end up with all function variables being global, which can cause subtle bugs. I **recommend** using `local` for any temporary variables inside functions. It

greatly reduces the risk of name collisions and unintended side effects. For example, in `giv_config()`, variables like `key`, `val`, `tmpfile` are currently global. If `giv_config` calls another function (or is called in a context where those vars were set), they could conflict. Making them local would encapsulate them. If you're worried about ShellCheck's POSIX warning, you can add a directive in `.shellcheckrc` to allow `local` (ShellCheck recognizes that pattern for dash). It looks like in `commands.sh` you added dummy initializations for some globals to pacify ShellCheck (like `__VERSION`, `GIV_TEMPLATE_DIR`, etc. set to dummy values) ³⁵. This suggests ShellCheck was complaining about use of undeclared variables. A cleaner way could be to declare them at top or via `local` when used. The dummy assignments trick works but might confuse future maintainers. Consider adding a comment there explaining they exist just to satisfy ShellCheck and have no functional effect (except setting default values which might actually mask bugs if those should be dynamically determined). Alternatively, use `# shellcheck disable=SC2154` for known external variables instead of assigning fake values. In any case, ensure these placeholders don't override real values. For instance, setting `GIV_OUTPUT_MODE="update"` by default might override what the user passed if not careful – hopefully by the time you use `GIV_OUTPUT_MODE`, it's already set to something else and not relying on that dummy. But if not, that line could cause a bug (imagine the user didn't specify `--output-mode`, do you set `GIV_OUTPUT_MODE` somewhere else or are you unintentionally defaulting it to "update" here?). If it's intentional as a default, fine – but be deliberate about it.

- **Quoting and Spacing:** I've scanned portions of the diff and see that you're quoting variable expansions appropriately almost everywhere (e.g. `"$GIV_CONFIG_FILE"`, `"$key"`, etc.). Keep that up. One place to double-check is constructing multi-line strings or commands. For example, when you assemble the OpenAI API JSON payload, ensure strings are properly quoted and escaped. Possibly you build a JSON by writing to a file or using `printf`. Watch out for `printf '%s' "$message"` where `$message` itself has quotes/newlines – though JSON requires quotes around content, so likely you use `jq -R` or similar or a heredoc approach.

- **Use of `test` vs `[]`:** Both are fine (they are the same `test` binary or builtin). Just be consistent. Also prefer `=` over `==` in single-bracket tests (I see you correctly use `=` ³⁸). And for numeric comparisons, use `-eq` etc. Standard stuff.

- **Pipelines and `set -e`:** Recall that in POSIX shells, `set -e` is tricky with pipelines – by default, `set -e` does **not** consider a pipeline as a whole failing if a non-last component fails. Some shells (bash with `pipefail`) can adjust that. Dash/ash do not have `pipefail`. So you may have cases like `grep -q something file | head -n1` where `grep` might exit 1 (not found) but the pipeline's exit status is that of `head` (which might be 0), thus `set -e` won't stop the script even though `grep` "failed". This can hide errors. In your code, I noticed you often check for conditions explicitly (e.g. using `if grep -q ...; then ... fi` which avoids that issue). Good. Just remain vigilant: any pipeline where an earlier command's failure matters should be refactored to an explicit check.

For example, in config

```
grep -v -E "^(key|givkey)=" "$GIV_CONFIG_FILE" | grep -v '^$' > "$tmpfile" 39
```

– here if the file doesn't exist, the first `grep` fails (exit 2), but the pipeline exit code will come from the second `grep`. If `set -e` is on and this pipeline is not part of an `if` or so, the script *might* continue even though nothing was done. It's minor because you guarded file existence above, so this pipeline should only run if file exists. Also, an empty config file would cause first `grep`

to exit 1 (no matches), but second grep still writes nothing and returns 1 as well (since no input, no match), thus pipeline exit status 1, which should trigger `set -e` to abort *unless* it's in a context that ignores it (in `mv "$tmpfile" "$GIV_CONFIG_FILE"` line, I think it's fine because you didn't combine with others). Just ensure any pipeline whose failure you want to ignore is either in an `if` or you append `|| true` explicitly (with a comment why). And any pipeline whose failure you *do* care about is either a single command or handled via a conditional.

- **Portability:** Also consider BusyBox environments (which cover many CI or container scenarios). BusyBox's `sh` is fairly complete but a few gotchas:

- `mktemp` exists in BusyBox, but older versions might not support long templates or certain options. Your `portable_mktemp` likely handles using `mktemp -t` vs something if needed.
- `sed -E` is **not** supported in BusyBox `sed` (BusyBox `sed` doesn't have `-E`, only `-r` for ERE, and in very old versions, no ERE support at all by default). In your config code, you use `grep -E` (BusyBox `grep` does support `-E`) and `sed -E` for substituting the key ⁴⁰. That `sed -E` will fail on BusyBox `sed`. To be truly POSIX portable, avoid `-E` in `sed`. You can usually rewrite the regex in basic syntax or use `grep -E` + `tr` combo as you did for other parts. In that specific case, you actually don't need `sed` at all to extract the key name; you could do it with shell parameter expansion as an alternative (as noted earlier). Since it's a simple prefix removal up to `=`, e.g.

```
key=${line#GIV_}
key=${key%*=*}
key=$(printf '%s' "$key" | tr 'A-Z_' 'a-z.')

```

This would replace the `sed` and two `tr` calls with maybe one extra step, but it's pure shell and avoids regex differences. Anyway, if you keep `sed`, use basic RE: the equivalent basic regex to `^GIV_([A-Z0-9_]+)=` would be `^GIV_[A-Z0-9_][A-Z0-9_]*`

`=` (since basic `grep/sed` requires escaping the group and quantifiers differently). That's uglier and not worth it here – simpler to avoid `sed`.

- Check other uses of `sed -E` or `grep -P` (hopefully none, since `-P` is not POSIX at all). Also ensure you don't rely on `echo -e` or `seq` or other non-POSIX utilities.
- One more: `head -n1` and `tail -n1` are fine (POSIX `head/tail` support `-n`). `read -r` is fine. `printf` is fine. `command -v` is POSIX (better than `which`). You seem to be doing well here.

Testing & Quality

Your tests under `tests/` are quite comprehensive, covering commands, metadata, version extraction, etc. This is great for ensuring the refactors didn't break functionality. A few suggestions on tests and maintenance:

- **Remove Obsolete Tests:** You already removed tests for the old providers (`test_provider_node_pkg.bats`, etc., gone in the commit [41](#)). Good. If any test references outdated environment variables or behaviors, update or drop them. E.g., if previously tests checked that `provider_python_toml.sh` outputs something, that should be replaced by tests for `get_metadata_value` logic.
- **Add Config Command Tests:** The commit [9e7b26e](#) added `tests/commands/config.bats` and updated `summary.bats` [42](#), which likely test setting a config value and then running `giv summary` to see if it picked up that config (maybe testing that an env var override works in summary context). Make sure you also test edge cases of `giv config`:
 - listing when file is empty vs non-empty,
 - getting a key that doesn't exist (should return non-zero and no output, I presume),
 - setting a key with special characters in value (e.g. a value containing spaces, or `&` or `#` – need to ensure those are handled, particularly `#` in a value would not be treated as comment in the middle of the line, which it shouldn't in shell as long as it's quoted in the config file).
 - unsetting a key removes it (and not others).
 - shorthand usage: `giv config some.key value` and `giv config some.key` produce the correct results.
- make sure a key with dot can be retrieved via both `giv config some.key` and `giv config --get some.key` for consistency.
- **Simulate Real Scenarios:** End-to-end style tests that mimic actual usage can be very valuable. For example, create a dummy git repo in `fixtures/` with a known version file and a couple of commits with TODO comments. Then run `giv changelog` on it and assert that the output contains the expected version and that TODOs were included. You likely have some of this (maybe `test_metadata.bats` and `summarize_target.bats` cover it). The more real-world, the better to catch integration issues between modules.
- **ShellCheck in CI:** If not already, integrate ShellCheck and maybe even a formatter like `shfmt` in the development process to keep style consistent. Given a `.shellcheckrc` is present, I assume you run it. Keep it up to enforce best practices as you refactor further.
- **Documentation vs Behavior:** Ensure that README and help messages are updated to reflect changes. For instance, the README snippet in this branch still shows `--model-mode auto/local/remote/none` and others, which is good. If any default behaviors changed (like no longer needing certain flags due to auto-detection), update docs. E.g., if you removed the need for specifying `--version-file` for common cases, mention that the tool auto-detects versions in standard files now. Also document the config file location and format (likely in `docs/configuration.md` as

referenced in the example config ⁴³). It's important for maintainability that code and docs stay in sync as the design simplifies.

Conclusion & Further Refactoring

The `spike/config` branch significantly improves the giv CLI's codebase by removing unnecessary parts and making the remaining code more modular and config-driven. To summarize the key refactoring recommendations:

- **Simplify and unify logic:** Continue to merge duplicate code paths (as done with project metadata). For any functionality that is implemented in more than one place, try to refactor it into a single function. For example, if both `summary` and `release-notes` need to loop over commits and format them, they could call a common helper to get a list of commits and then each apply their own template.
- **Lean on POSIX shell strengths:** Use shell builtins for string manipulation where feasible (prefix/suffix removal, parameter expansion) to avoid external calls. This makes the code faster and more portable. We identified one such case in config key parsing where shell string ops can replace `sed`.
- **Robust error handling:** Make sure every external call (git, curl, jq, etc.) is checked. The user should get a clear message if something fails (e.g., "Git not installed or not in PATH" if `git` command fails – though given the audience that's unlikely, but still).
- **Maintain POSIX compliance:** Avoid bashisms (we didn't spot any obvious ones except maybe the use of `sed -E`). Test the CLI on minimal environments (Dash shell on Debian, BusyBox shell, etc.) to ensure nothing breaks. You might add a CI job that runs tests with `/bin/dash` as the interpreter (e.g., run bats with `SHELL=/bin/dash` so that each test uses dash internally – Bats may require bash itself to run, but you can at least have giv executed by dash within the tests).
- **Keep it familiar:** The more your commands behave like analogous git commands, the easier for users. The config command is a great example. Also consider naming and output conventions: e.g. `giv available-releases` lists script versions (like `git tag` or `git releases`). It should format them clearly and maybe indicate which is current. If not done, maybe sort versions according to semver (non-trivial in shell, but you could pipe to `sort -V` if GNU sort is available, or use a small version compare function). It's those little touches that make the tool feel polished.

In terms of architecture, the current breakdown (with distinct scripts for separate concerns) is already quite good for maintainability. New contributors can work on one area (say, improving `history.sh` logic) in isolation without touching the LLM or config code. Just ensure the interfaces between these modules are well-defined and documented (for example: what does `get_metadata_value` expect as input and return as output? Document that it echoes the value or nothing and returns non-zero on failure, etc.). Comments in the code explaining non-obvious tricks (like the dummy variable assignments for ShellCheck) will help future maintainers understand why things are done a certain way.

Overall, the branch is moving in the right direction: **simpler, more flexible, and more familiar**. By addressing the minor issues above (regex escaping, portability, fully removing obsolete code, and tightening up a few logic paths), you'll have a robust codebase that's easy to maintain and extend. Keep using ShellCheck and your test suite as you refactor – they will catch many potential mistakes. And as always with shell scripts, consider the user's perspective: make error messages clear, avoid surprising behaviors, and document any environment variables or config knobs that control the behavior.

With these refinements, **giv** will remain a zero-dependency, developer-friendly tool that cleanly integrates with Git workflows to generate useful AI-assisted content.

1 2 3 4 5 12 25 26 27 28 29 31 32 33 34 35 36 37 41 43 Refactor project metadata handling and improve environment loading · giv-cli/giv@f1bd745 · GitHub
<https://github.com/giv-cli/giv/commit/f1bd745b124c2d5bc2e5cc151e1c55ff9b230c0a>

6 7 8 9 10 11 13 14 15 16 17 18 19 20 21 22 23 24 30 38 39 40 42 Refactor config command to enhance error handling and add support for... · giv-cli/giv@9e7b26e · GitHub
<https://github.com/giv-cli/giv/commit/9e7b26ef9ff65a6bec6e5c505c330db230f76d0c>